Secure Local-Network-Only Retrieval-Augmented Generation

Presented by,
Mridula J.B
Neeraj Ram B
Muhammed Nabeel

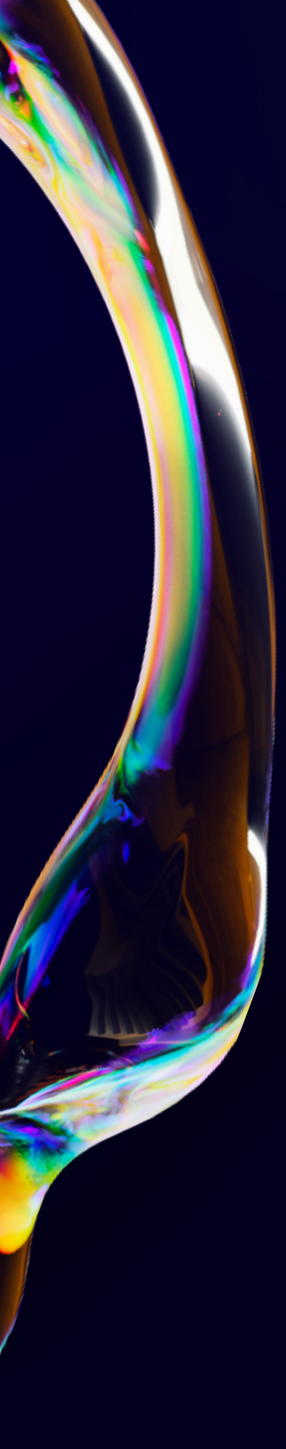

Table of contents

1.	Introduction
2.	Literature Survey
3.	Research Gap
4.	Problem Statement
5.	Objectives
6.	Design
7.	Conclusion
8.	References

Introduction

- This project focuses on developing a secure, local-network-only Retrieval-Augmented Generation (RAG) system for managing employees, projects, and sensitive internal documents in an IT organization.
- The system operates entirely offline within a local network, ensuring data privacy, controlled access, and zero internet dependency.
- By combining role-based access control with intelligent document retrieval using a local AI model, the system enables secure storage, organization, and querying of internal data.
- The project aims to provide a reliable, efficient, and privacy-focused internal management solution for enterprises.

Literature Survey

Paper 1: "Startups’ roads to failure."

In the article "Startups’ Roads to Failure" (published in Sustainability, 2018), Marco Cantamessa and his colleagues provide a systematic analysis of why startups fail. While most entrepreneurship research focuses on success stories, this study examines the "post-mortems" of failed ventures to identify recurring patterns.

Core Methodology: The SHELL Model

The researchers adapted the SHELL model—a framework originally developed in the aviation industry to analyze aircraft accidents—and applied it to the entrepreneurial context. The acronym stands for:

S (Software): The business model, processes, and intangible rules.

H (Hardware): The physical product or the technology being developed.

E (Environment): The external market, competition, and regulatory context.

L (Liveware - Organization): The internal team, founders, and organizational dynamics.

L (Liveware - Customer): The relationship with users and the market.

Paper 1: "Startups’ roads to failure."

Key Findings

The study analyzed a database of 214 startup post-mortem reports from Autopsy.io. Their primary conclusions include:

1. Dominance of Business Development Issues: The single most common reason for failure is the lack of a structured Business Development strategy. Many startups focus too heavily on the product (Hardware) while neglecting how to build a scalable and sustainable business around it.
2. Product-Market Fit (PMF): A massive number of startups fail because they build products that the market does not actually need.

- Timing of Failure:Under 1 Year: Failures are often due to inexperienced management and a total lack of a business model.
- 2–3 Years: This is the "critical window" where most failures occur. Startups often run out of cash here because their initial business model fails to generate traction or attract further investment.
- 4–5 Years: Failures at this stage are frequently caused by internal friction or misalignment between founders and investors regarding the company's long-term goals.

3.Financials are a Symptom, Not a Cause: While "running out of cash" is a frequent reason cited by founders, the researchers argue this is usually a consequence of deeper issues like a flawed business model or poor marketing, rather than the root cause itself.

Paper 2: Retrieval-augmented generation for knowledge-intensive nlp tasks.

The study demonstrates several advantages of the RAG approach over traditional "closed-book" models:

- Reduced Hallucinations: By forcing the model to retrieve specific document fragments before generating a response, the system significantly improves accuracy and trustworthiness.
- Knowledge-Intensive Performance: RAG out-performs standard models on tasks that require precise factual knowledge, such as answering questions about niche organizational data.
- Dynamic Updating: Unlike traditional models that require expensive retraining to learn new facts, a RAG system can be updated simply by adding new documents to the external index

Paper 3: What is retrieval-augmented generation

Core Concepts & Analogies

- The "Open-Book" Analogy: Standard LLMs answer from "memory," which leads to errors. RAG acts like an "open-book" exam, querying a library of trusted documents for factual accuracy.
- Two-Step Pipeline: A "retriever" first pulls relevant data from a document index, then a "generator" (LLM) uses that specific data to craft an answer.

Strategic Organizational Benefits

- Trust & Transparency: Provides a trail of sources so users can verify AI claims.
- Cost-Efficiency: Updates the system by refreshing the document index rather than performing expensive model retraining.
- Real-Time Accuracy: Accesses the latest files, bypassing the knowledge "cutoff dates" of traditional LLMs.

RESEARCH GAP

1. **Cloud-Dependent AI Systems**

Most AI document assistants rely on cloud processing and require internet connectivity.
Gap: No sufficient solutions that allow AI document querying completely inside a private local network.

2. Data Privacy in Document Handling

Existing document management platforms store and process sensitive company data on external servers.
Gap: Lack of secure AI-based querying of confidential internal documents without data leaving the organization.

RESEARCH GAP

3. RAG Without Access Control

Current RAG systems mainly focus on improving answer accuracy.
Gap: Limited research integrating RAG with role-based access control for different organizational users.

4. Fragmented Organizational Tools

Employee management, project tracking, and document storage systems operate separately.
Gap: Absence of a unified platform combining management functions with intelligent AI retrieval.

RESEARCH GAP

5. Offline AI Deployment

Most RAG implementations depend on online models and external APIs.
Gap: Very few systems implement a fully offline, local-network AI model for private enterprise environments

Problem Statement

Our project addresses the challenges of insecure storage of sensitive organizational data, inefficient information retrieval due to limited intelligent search capabilities, and fragmented management systems by proposing a unified solution that enhances data security, enables intelligent and accurate document retrieval, and integrates employee, project, and document management.

Objectives

01

Secure and Private Architecture

- Develop a production-grade RAG system operating exclusively on a local IP network to ensure total data sovereignty and zero internet dependency.

02

Intelligent Knowledge Retrieval

- Implement a local vector search engine and dense document index to provide up-to-date internal data, minimizing hallucinations and improving response accuracy.

03

Role-Based Access Control (RBAC)

- Integrate a secure authentication framework with tailored permissions for organizational roles ranging from Employee to Founder

04

Unified Enterprise Management

Consolidate project management, secure messaging, and activity logging into a single platform to mitigate operational risks and inefficiencies.

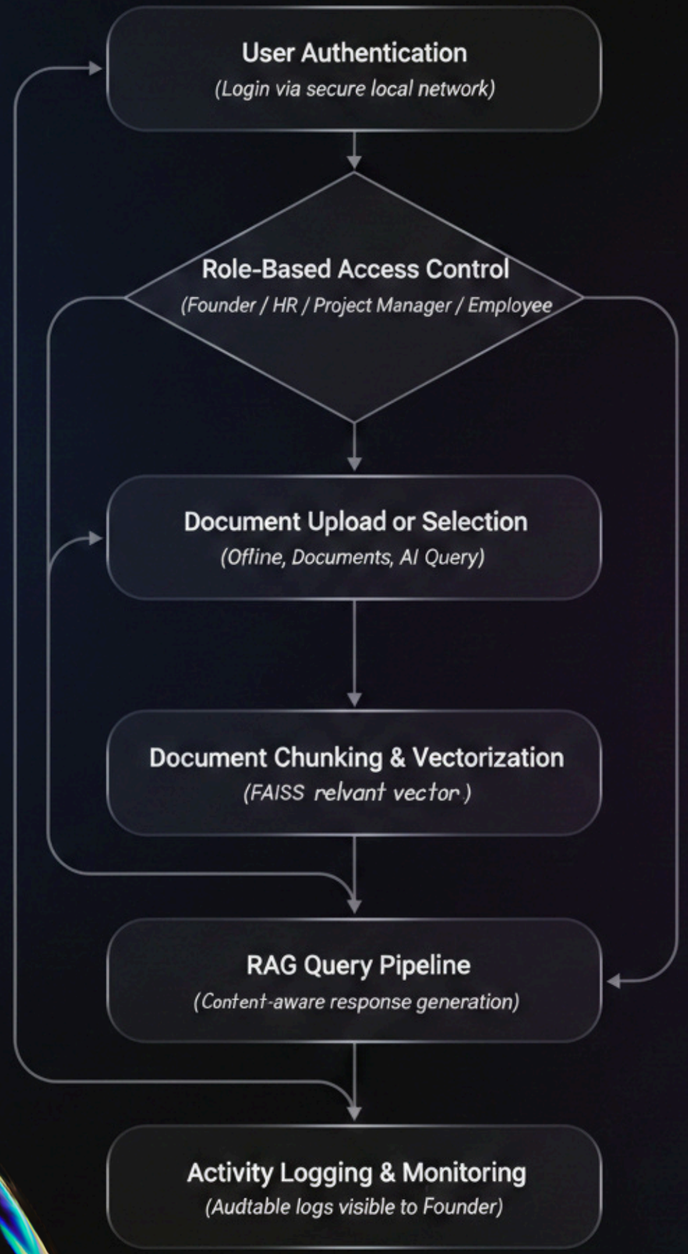

Design

- **Secure Local Access:**

  The system operates entirely within a local network, ensuring data privacy through authentication and role-based access control.

- **Integrated Project & Document Management:**

  Projects and documents are securely organized with controlled permissions for different user roles.

- **Intelligent RAG Pipeline:**

  Documents are chunked and vectorized using FAISS to enable accurate, context-aware AI responses via a local LLM.

- **Monitoring & Audit Logs:**

  All user activities are logged to provide transparency, accountability, and administrative oversight.

CONCLUSION

1. **Secure Local Operation**
   - The proposed system operates completely within a private local network, ensuring full data privacy and eliminating internet dependency.
2. **Intelligent Document Retrieval**
   - The RAG-based pipeline allows users to query internal documents using natural language and obtain context-aware information.
3. **Role-Based Access Control**
   - The system integrates authentication and permission levels for different organizational roles, restricting access to authorized data only.
4. **Unified Management Platform**
   - Employee records, projects, and documents are managed within a single system, reducing fragmentation and improving efficiency.
5. **Practical Enterprise Solution**
   - The platform combines local AI processing, vector search, and activity monitoring to provide a reliable and privacy-focused internal management solution

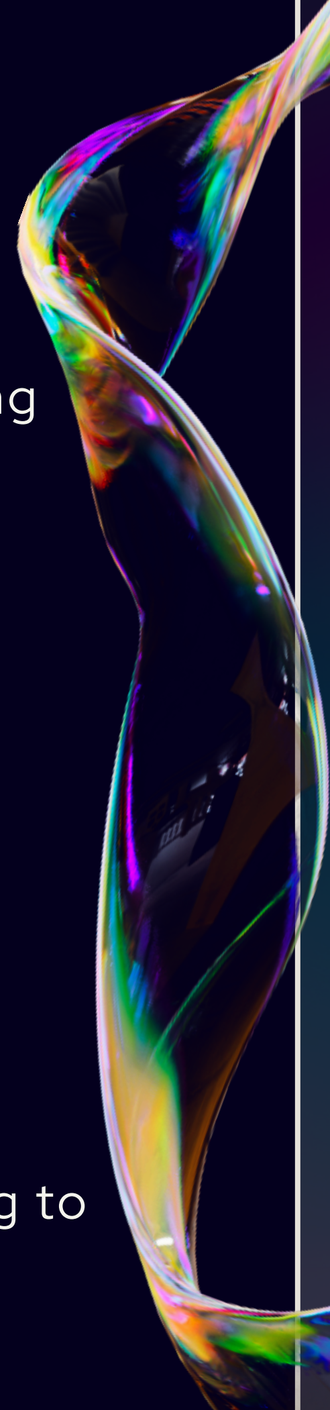

REFERENCES

[1] Marco Cantamessa, Valentina Gatteschi, Guido Perboli, and Mari angela Rosano. Startups' roads to failure. Sustainability, 10(7):2346, 2018.

[2] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich K¨uttler, Mike Lewis, Wen-tau Yih, Tim Rockt¨aschel, et al. Retrieval-augmented genera tion for knowledge-intensive nlp tasks. Advances in neural informa tion processing systems, 33:9459–9474, 2020.

[3] Kim Martineau. What is retrieval-augmented generation (rag). IBM Research, 2021.

Thank You